

# Running tasks in parallel with Celery

Irina Colgiu

Human Genetics Informatics

Wellcome Trust Sanger Institute

# What is Celery?

**Celery** is a task management and queueing system

# Celery vocabulary

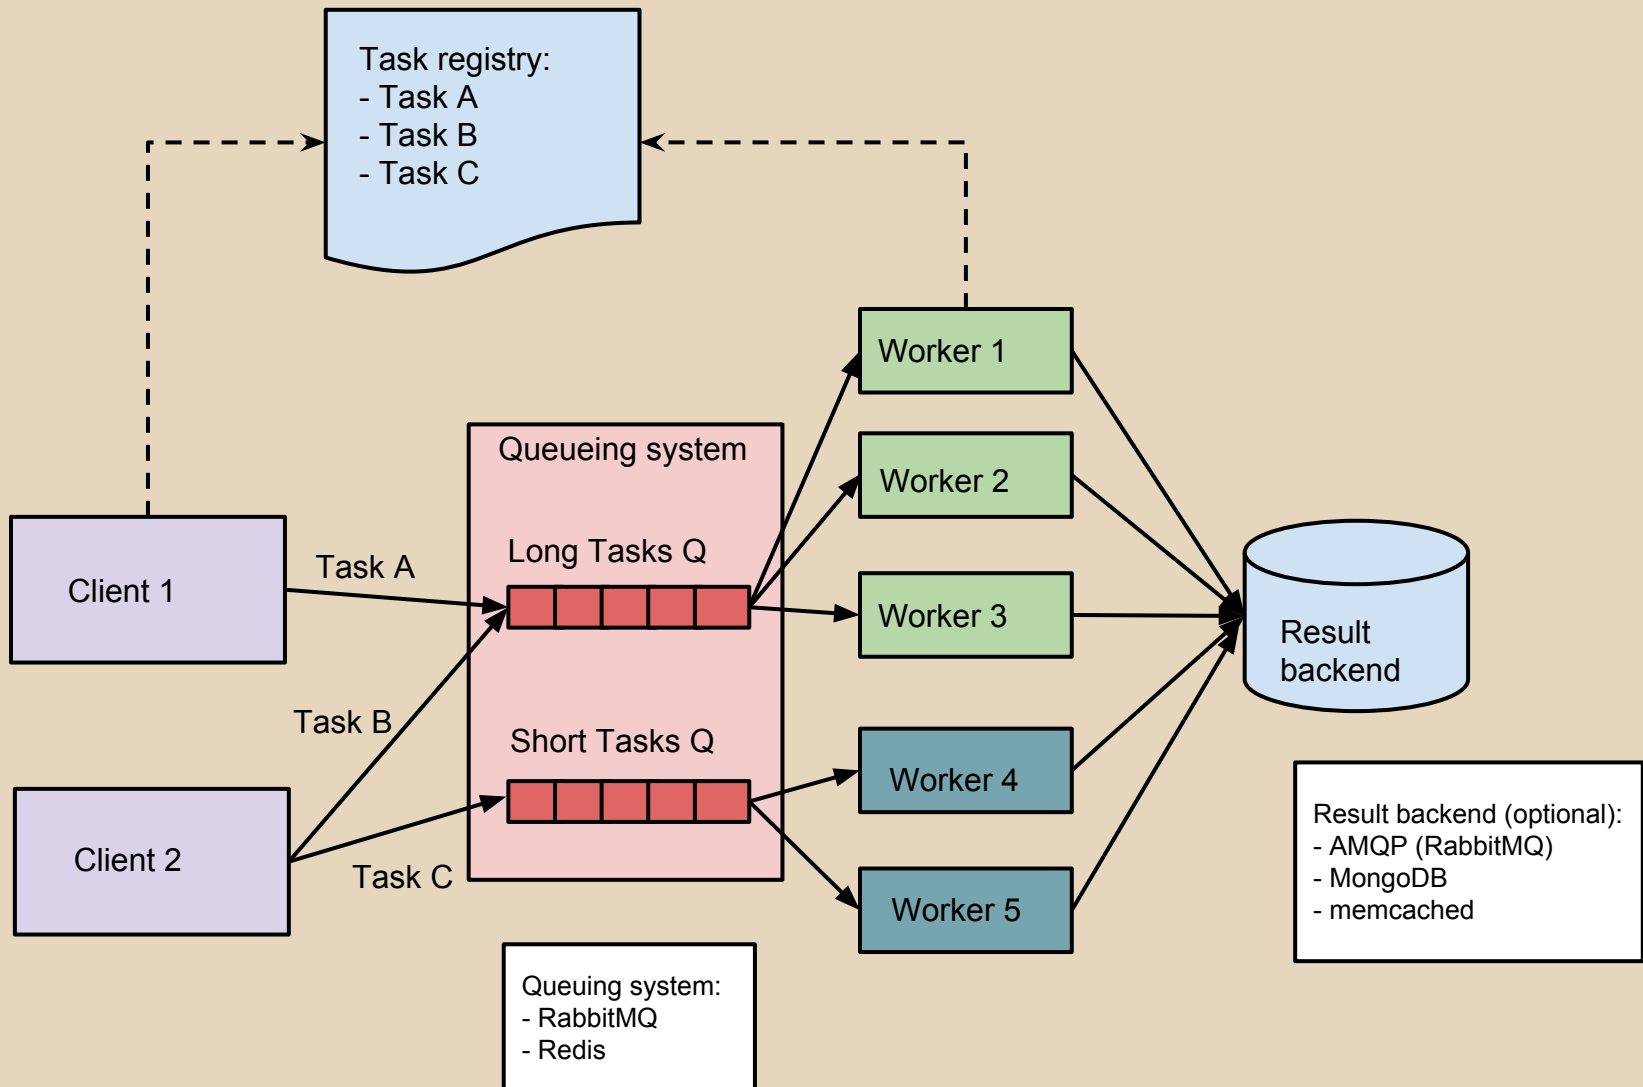
Task = a unit of work (a function defined by the user)

Worker = a process that executes tasks

# What Celery offers

- A solution for executing work **async** (or **sync**) on **one/more machines**, as **one/more processes/threads** **remotely**
- Managing task dependencies
- Dealing with task failure
- Handling machine failure
- Task execution updates (useful for long-running tasks)
- Collect results from tasks
- Monitoring the tasks

# Celery functionality



# How to execute a task?

Defining a task:

```
broker = 'amqp://guest@machineA:5672'
celery = Celery('tasks', broker=broker)

@celery.task
def add(x, y):
    return x + y
```

Defining a results backend:

```
...
celery = Celery('tasks',
                broker=broker,
                backend='amqp')
...
```

Calling a task:

```
# ASYNC - the task executed in a worker
# process

async_result = add.apply_async((2, 2))
```

Collecting the results:

```
async_result.ready()
# True/False

async_result.state
# PENDING/STARTED/RETRY/SUCCESS/..

async_result.result
# Contains the return value of your task
```

# How to start a worker?

```
$ celery worker -Q LongTasksQ
```

# Useful options

- **Task routing** - one/more queues, the workers can listen to one/more queues
- **Task timeout:**
  - **soft limit** - the task is considered as failed if it doesn't finish before the timeout => requeued
  - **hard limit** - restart the worker



# Useful options

- **Task retry** - configure the max retries, delay between retries, delay before the first retry
- **Rate limit** - limiting the nr of tasks to be executed in a given time frame
- **Autoscaling** - resize the worker pool depending on the load

```
$ celery worker --autoscale=3,10 -Q LongTasksQ
```

# Limitations

- The tasks **MUST** be predefined - in the *tasks registry* by the time you start the worker processes
- You need to start and kill the workers “by hand” (unless you work in the daemon mode)
- Each type of results backend has its own strengths and weaknesses - e.g.: amqp (queues) - each task result is reported on its own queue

# So what do I need to install?

```
$ pip install Celery
```

Either of:

- RabbitMQ
- Redis
- Database (MongoDB)

as message broker & results backed.

# Where to read more

<http://www.celeryproject.org/>

**Thank you for your attention!**

**Questions?**